

Project: CSE 241 Database Systems

Fall 2025

Goal:

This project aims to provide a realistic experience in the conceptual design, logical design, implementation, operation, and maintenance of a relational database and associated applications. First, we shall describe the application, then the categories of requirements, and then some suggestions on how deeply you need to go in each category. A real project would require a substantial development team working for several months (or more). You will do this alone over several weeks. We have chosen to go with individual rather than group projects because the goal of this project is for you to gain a personal appreciation of the depth and breadth of issues that go into the design of a database application rather than have you specialize in just one aspect (and rely on others for the rest).

The project can go well beyond the minimal requirements. We encourage such extensions but ensure your project will run in our testing environment.

The description of the enterprise you are modeling is somewhat vague and incomplete. This is by design — in real life, *your* “customers” are managers in the enterprise who are not computer science experts. You must fill in the holes in this document to create a precise design and concrete implementation of the interfaces and applications using the database.

The checkpoints specified for the project are designed to help you get some feedback along the way and keep you on schedule. Please pay particular attention to the requirements for project submission and the testing protocol we shall use. Because of the large number of projects that need to be evaluated in a short time, the testing protocol will be strict, and exceptions will not be granted.

Enterprise description:

The enterprise is *LUShop*, a fictional online marketplace. Customers connect to the website with credentials and can browse and purchase items. For this project, we will focus solely on operations that impact the financial aspects of LUShop. Hence, we shall concentrate on customers, managers, reporting, financing, and items/services to be purchased. LUShop is an online service without physical locations.

LUShop offers a variety of operations:

- Accounts: Customers and Managers are users of LUShop. Their names are their access credentials (no passwords). Customers can purchase items, apply for financing, and browse the catalog. Managers populate the catalog and produce reports. Customers can be individuals or businesses. An individual can purchase items, while a business can only buy services.
- Catalog: LUShop offers a variety of items and services that customers can purchase. These items/services are maintained by LUShop’s managers. Items have an identifier, a searchable description, a price, and a vendor. Services also have a duration.
- Financing: Services can only be paid through a bank account, while items can be purchased using credit cards or installments. When a customer uses installments, different plans with different terms and conditions are shown. Managers are responsible for populating these installment plans.

Interfaces:

Before your “real” interface runs, prompt the user to enter the Oracle password for your Oracle account. This keeps your password out of your Java source code. We shall be running a script at the project deadline to change your password to something we know, so having your password in your Java code is insecure and ensures that your final project will not execute for us when we evaluate it. DON’T DO IT!

You should implement the following interfaces:

- Customers. This interface handles customers’ information, payment methods, and records all purchases and total expenses.
- Catalog. This interface allows customers to browse the catalog and purchase items/services.
- Manager. This interface allows managers to produce at least three aggregated reports.

You should implement your interfaces in one executable with an initial dialog, allowing the tester (that is, a member of the course instructional team) to choose an interface.

Each interface should include features to allow users to find the needed information. For example, when a customer purchases an item, the list of credit cards on file is displayed rather than expecting the tester to remember such information.

You should implement ALL the listed interfaces. There are a lot of commonalities of features among the various interfaces both in terms of the actions to be taken and the conditions to be tested. Take these into account and modularize your code (in Java and via database functions, procedures, or triggers) so that each individual interface has only a modest amount of interface-specific code.

Your README file should tell us about each interface. Each interface should be designed to be usable by the customers or employees of the enterprise and not contain SQL jargon in its dialog.

- **Code quality:**

Just as you would not expect software that you buy or use to crash, throw exceptions, or fail to handle input errors gracefully, we expect well-written, error-tolerant code. Interactive interfaces should test for proper input (data type, range, and validity) and allow convenient re-entry after an error.

You should not expect your users to have the database memorized. It should be possible for users to query for information.

Your code should *never* throw an exception. Catch exceptions, print out a reasonable message, and have the code carry on reasonably, or, if no reasonable continuation exists, terminate gracefully with a message to the user.

- **Testing:** Test your code in Sunlab. Your code might run on your machine but fail in Sunlab. Your code must work in Sublab, because there is where it will be tested when graded.

- **Interface details**

As we note again later, projects will be tested only *from the command line on the sunlab machines* and possibly via remote login. This means that not only do you not get credit for having a GUI, but you will perhaps lose all credit for the executable since GUIs do not run over a remote terminal window. We realize that a command-line interface is a last-millennium concept, but it simplifies the overall project and provides us with major efficiency gains when we have to evaluate all projects in a tight time window.

- For whatever interfaces you write, you may use as much PL/SQL as you wish. We generally prefer (but do not require) that you have the database system do the work via PL/SQL rather than your non-JDBC Java code do much of the work. Making effective use of what Oracle has to offer can compensate for a lost point of two elsewhere.

What you need to do:

There are several steps to this project. Although you will inevitably need to go back and change things as you move along, it is desirable to do a very good job at each step to reduce the amount of work that winds up being redone. Note that “very good” does not mean “perfect.” If you take too long to make the early stages of the project perfect, you’ll find yourself pressed for time at the end. You’ll need to stay on schedule and learn, as one has to in real life, when to “declare victory and move on.” If you manage your coursework on an “earliest deadline first” basis, you’ll find yourself in deep trouble on this project. But if you stay ahead of our suggested pace, you’ll likely have a quality project.

Here is a set of stages to follow:

1. **ER design** For this project, you are **required** to use the primary notation used in the textbook and not any of the other notations that appear at the end of Chapter 6, nor any other notations that you may find online. Note also that ER diagrams are *not* schema diagrams, which appear in Section 2.4 of the text and which we did not cover in the course.

Construct a good, complete ER design for the enterprise. There will be a formal checkpoint in which a grader or TA will review your design. It is worthwhile to refine this design to represent the enterprise with a considerable level of detail. That makes the next steps much easier.

To start, it is best to sketch the diagram with pencil and paper. Not only are there a lot of changes initially, but often you discover that the placement of entity sets on the page can influence how many lines cross. Relocating entity sets physically on the page can clean up a diagram considerably.

Note that a good ER design includes careful choice of what things are entity sets and which are relationship sets, proper placement of attributes, use of generalization/specialization where appropriate, etc. A common error is to think relationally and then reverse-engineer the ER design. That approach often leads to hidden relationship sets encoded in common attributes between entity sets. Foreign keys are a relational concept; they are not a feature of ER designs and such errors are one of the prime ways students lose points on the ER section of the grading rubric.

As you make decisions, include notes explaining the assumptions you made about the enterprise leading to those decisions. That will help you when you go back and reconsider your design. You may want to turn in some of those notes with your diagram to help us understand how you view this enterprise.

Once your design is well along the way, you will need to create an electronic version. If you use software tools for this, be sure the tool can generate a pdf file. **We shall be accepting only pdf for ER diagrams.** If you hand-write your diagram, be careful that your writing is legible and be careful that your scan/photo is also readable.

2. **Relational schema** The textbook gives a set of rules for generating a relational schema, including primary-key and foreign-key constraints directly from the ER design. If your ER design is good, you will be nearly done at this point. There may be some data dependencies that were not captured in the ER design that may lead to some further normalization. Check for this, but for a good ER design, there won’t be many, if any. You may decide to add some additional indices for performance. You may decide to add some triggers or stored procedures later on. Those don’t all have to be done at the start of your work; you can always add them later.

Once you have a conceptual version of your relational database schema, you need to generate a SQL DDL version of it. That means deciding on reasonable data types and getting the syntactic details

of SQL right. Enter this in Oracle under your account. By default, only you (and the instructional team¹) have access to these tables.

Note that if at any point you find flaws in your database design, you need not only to fix the design in Oracle but also make any changes that may be required in your ER design. When you submit the final project, your ER diagram should be consistent with the database you have created in Oracle.

Hint: Don't type your DDL directly into SQL Developer. Instead create a plain text file with your DDL and copy/paste it into SQL Developer. If there are errors, edit the file and then copy/paste again. This allows you to retain a full version of your DDL for future editing without having to extract it from Oracle. This is important for several reasons: (1) should we have a catastrophic failure of the system, you can restore your schema quickly on a replacement machine (2) If you take a look at the SQL DDL Oracle generates for your schemas, you'll see that it is unnecessarily complex since it includes specification of all sorts of parameters for which we are using the default.

Note that you have to drop tables before you can re-create them. Just put the drop statements at the start of the DDL text file. Also note that foreign keys can't reference a relation that does not exist. So be sure to list the **create table** commands in an order that ensures that referenced tables are created before referencing tables. To drop tables, you need to drop the referencing tables before the referenced tables.

3. **Data generation and population of relations:** You need to put data into your tables. Include enough data to make answers to your queries interesting and nontrivial for test purposes, but there is no need to create huge databases. We have set a storage quota that is more than enough for any reasonable course project.

To avoid a typing marathon for data generation, write a program to generate test data or use data you can find online (but use only data that may be copied legally).² You can get some data fairly easily without much typing. You should automate data loading so that if you need to redesign part of your database or your interface code trashes the data due to a bug, you can reload your data without too much effort. A trick in the automated process is working around referential integrity constraints. Inserts have to be done in the right order to avoid a foreign-key violation.

Note in your README file.

There are a couple of situations that merit special attention:

- **Representing time:**

The format of **timestamp** data is detailed, and it is easy to make mistakes. Thus, you must be really careful when testing user input of time data. For good quality output, you likely will need the *to_char* function to get the format of output you'd like.

- **Generation of unique id values**

Creating unique identifiers automatically in Oracle took a bit of work until the introduction of the autoincrement feature in Oracle 12. Details appear in Section 4.5.6, page 160 of the textbook.

- **Bulk loading**

Oracle has a proprietary bulk loader whose use we do not cover in this course, but simply generating insert statements and running them is probably easiest.

4. **Interface coding:**

Don't forget the basics of good programming.

¹So, don't keep any personal data in your tables.

²Many prior students have used the mockaroo website to generate data, but regardless of how you generate data, you need to ensure that not only is each table reasonable but also that the linkages between tables (foreign keys) are reasonable. In this, as in other cases, investing in writing code to automate your work is most efficient.

- Check user input for being valid. If you are inputting an `int` using `nextInt`, first check that the user did not enter a nonnumeric character (recall `hasNextInt`). Produce good-looking output. And so on. Then when the user does something wrong, don't just quit. Provide a chance to try again without requiring unnecessary re-entry of input data.
You may note that our sample code in lecture does not do this. That's only because we're trying there to focus on the new material and avoid code that distracts from our main point. But for the project, we expect better quality code.
- Make good use of classes and methods in your Java code. Use PL/SQL triggers, functions, and procedures where appropriate. Although the grading rubric is about database functionality and not Java coding, we will deduct points for seriously poor coding and perhaps give a point or two extra for excellent use of database procedures that help streamline the Java code and ensure database integrity.
- Integrity checking: A well-designed database will protect against many types of bad updates. But others may not be easy to express using SQL constraints. Think about bad things we might enter using your interfaces and try to protect against them in your interface code if your database integrity constraints and triggers are not enough to do the job. In all cases, avoid having your code crash on an exception. Catch them and do something reasonable. That includes catching exceptions thrown by JDBC methods due to errors generated by Oracle. Note that we may attempt SQL-injection techniques to try to cause errors or exceptions.
- Concurrency: In real life, lots of updates and queries would be run on this database concurrently. We should be able to run several instances of your code from separate terminal windows and not run into anomalies. In most cases, Oracle's default concurrency will do well enough, but see the note below about self-inflicted concurrency disasters.
- Test with care. Get the basics running first. Unit testing is your friend.

5. Self-Inflicted Concurrency Disasters and Zombie Attacks:

Although real-life concurrency is not likely to cause much trouble in this project, you can get yourself into trouble as you test your code.

If your Java code with JDBC calls to the database terminates abnormally (i.e., it throws an exception) or you simply forget to close your connection, your code may leave behind zombie transactions that Oracle thinks are still active. If those transactions hold locks, your subsequent test runs may wind up waiting for those zombie transactions to complete. You'll see this as the system "hanging" without explanation. Eventually, in a matter of minutes, the connection (and its transactions) will time out and all will be well again.

You can avoid this by careful testing. Test your plain Java code before you add JDBC code. Be sure to catch every exception and then close everything before your code terminates. The try-with-resources construct in Java is a big help here. Exceptions can occur anywhere: during connection, during SQL execution, during result fetching, and during execution of plain Java code (array bounds, for example).

In principle, the DBA can kill transactions manually. More likely, by the time we'd get to this, the connection would have timed out anyway. For this reason, plus the large number of projects, we don't offer transaction-killing service and ask instead that you be patient with the timeout period if something slips past your attempts to code carefully.

An even worse condition occurs if you do not close your connections properly, Oracle sessions may survive for quite some time even though you kill the operating system process that ran your code. These zombie sessions, if sufficiently numerous, impede the function of the system. We can tell who is leaving such sessions behind, and while we realize the occasional zombie is inevitable, students who generate large numbers of zombies may get points deducted even before grading begins since this is evidence of bad testing methodology.

6. Password Security:

Your Oracle password is needed to run JDBC, but you are **forbidden** to include your password in plain text in your code. Instead, prompt for the user to input the password at the start of each of the assigned programs. We don't need to know your password. For grading, we shall use DBA authority to change your password and we can enter that new password to test your code. *If you ignore this and you do hard-code your password, your code will no longer work once we change your password, resulting in your code failing to execute.*

Checkpoints: There are several checkpoints scheduled. These are set in order to keep you on target for completing the project on time.

- **Checkpoint 1: September 25.** The one point for the project allocated to this checkpoint will be awarded for a *reasonable* ER diagram. It need not be perfect to get the point (and there is no partial credit).

For this project, you are required to use the primary notation used in the text and not any of the other notations that appear at the end of Chapter 7, nor any other notations that you may find online. Note also that ER diagrams are *not* schema diagrams, which appear in Section 2.4 of the text and which we did not cover in the course.

The ER grading will be much more stringent in the final version submitted at the end of the project. Thus, a perfect score on this checkpoint is not a guarantee of a perfect score for the ER component on the final version of the project; it is just an indication that you are on pace at this point.

The graders will provide feedback on your submission, and you are encouraged to attend office hours to check your ER revision.

- **Checkpoint 2: October 2.** At this point, you should have your relation schemas created in Oracle. At some point after this date, one of the graders will look online to see that they are there and include reasonable key declarations (including foreign keys). We don't expect to see "fancy" features like triggers or stored procedures, but it is certainly fine if they are there. The whole point of this checkpoint is simply to keep you on pace. The one point for the project allocated to this checkpoint will be awarded for a reasonable set of relation schemas. As for checkpoint 1, it does not need to be perfect. We will take a more careful look at your schemas when we evaluate the final project. So, as is the case for the first checkpoint, a perfect score on the checkpoint is not a guarantee of a perfect score on the relational design component of the final version of the project.

At this point, you should also have a plan in place for user-interface development. We shall not be reviewing that plan at this time.

- **Checkpoint 3: October 10.** We shall check online to see that your relations have been populated with data. We shall award one point allocated for this checkpoint if most of your relations have a reasonable amount of data.

You should have some degree of user-interface functionality at this point, but we will not check that.

- **Project due dates.** The project will be due in 3 parts: the ER design, the relational data, and the executable project code, each with its own due date.

– **ER design:** *Wednesday, December 3 at 11:55 PM hard deadline.* Submit on CourseSite a single pdf file with your ER diagram (just one file, not several zipped ones, no formats other than pdf). CourseSite will enforce the deadline. ER designs not in the required format may receive reduced or zero credit.

– **Relational data & Executable code:** *Friday December 5 at 11:55 PM hard deadline.* We evaluate your relational design and data in the tables, along with testing your executable code.

Read the details below on what to turn in very carefully.

What to turn in with your executable code:

1. The top-level directory should be named using your loginID and last name (for example, lbj265johnson). This top-level directory should then be compressed using the zip format under the same name (for example if the folder is lbj265johnson, the zip file should be lbj265johnson.zip). No other compression formats will be accepted for this project. To zip a project on the sunlabs, use the command “zip -r ZipFileName FolderName”.

Submissions that use any format other than zip will not be read and will receive a score of zero. (So no *rar* format files and no tarballs)

2. A README file at the top-level in the directory hierarchy that explains what is where, etc. Include usage instructions for the interfaces (perhaps suggestions of good customers to test, for example). Also include sources of all data and code obtained from others (note the collaboration rules below). Be sure to include your name at the start of the README file. If you have done anything extra, be sure to point it out here!

The goal here is to ensure that we see all the goodness in your project and to help us have a good first experience with your code. A project that opens with “enter your ID” is not very friendly to a grader who has no idea what ID to enter. A good README file can provide some first suggestions. We’ll read your database to find others and create our own new users too.

Your README file must be in PLAIN TEXT – not doc, pdf, rtf, or anything else. It must be named either README or README.txt, but note that having a .txt extension is not enough to guarantee plain text format. It must be possible for us to run “cat README” in a terminal window and read your file without having to transfer it to our own machine so as to be able to open it with some app. Test your file to be sure this is the case.

3. One executable jar file that provides access to all your interfaces. The file is to be named xyz123.jar, where xyz123 is your Lehigh loginID. We’ll run “java -jar xyz123.jar” in a terminal window on a Sunlab machine of our choice, so you cannot rely on special settings for environment variables, etc. You cannot rely on us having the Oracle ojdbc jar file or oracle subdirectory in any specific place, so put it in the right spot yourself. If your jar file fails to run, you get a zero for the executable part of the project. Since you have the ability to test this yourself, there is no reason for us to attempt to do debugging.
4. One directory containing your Java project code, named xyz123 where xyz123 is your Lehigh loginID. If we wish, we may recompile your code. We’ll do this either by “javac *.java” or use a makefile you provide. If there is a makefile to use, it must be noted in the README file. Note that we will compile using the default Java version in the Sunlab and will NOT do it in the framework of any specific IDE. If we decide to compile your code and compilation fails you may get a zero for the executable part of the project.
5. A second directory for any data-generation code you wrote. We most likely won’t test this, but your code needs to be there for our review.
6. We will use DBA rights on Oracle to change your password and thus we shall be able to look at your tables. This means that your data must be on the Oracle system we are using for our course and cannot be on a personal installation of a database system. Our tests may modify your database and we will not restore your database to its status prior to our tests. (Also this means you should not be using your course account on Oracle for any personal data or for any other course, since you’ll lose access to such data, but we will have access!).
7. Do NOT turn in a listing of all your data. We can see them online.

Grading: We shall use the following approximate template for grading:

1. Checkpoints:
3 points, 1 each for Checkpoints 1, 2, and 3.
2. ER design: 6 points
3. Relational design, including constraints, triggers, and indices: 6 points
4. Data creation: sufficient quantity, reasonable realism, sufficiently “interesting”: 3 points
5. User interfaces, including proper features, proper updating of the database, etc.: 15 points

Note that bad database design can lead to interface problems, so design issues actually can have a larger impact than just the points specifically assigned to them.

Getting details right is important. Something as minor as throwing an exception on user input can cost a couple of points. Just like when you buy software yourself, the evaluation is not based on the percentage of code that is correct! We'll be evaluating your code as if we were users/managers at BRC.

6. We reserve the right to give extra points for exceptional solutions to parts of the project. We also reserve the right to deduct points in the unlikely event that we identify problems not covered by the items above.

7. **Draft submissions:** CourseSite will be set up to allow you to upload a draft without clicking submit. We won't grade drafts until after the deadline. But if you click submit before the deadline, that constitutes a submission, and we will then be free to grade your submission before the deadline at our convenience.

Submitting a draft is a good way to protect yourself by submitting something that mostly works while you continue to make improvements. That way, you have something there just in case your fixes turn out to fail badly.

8. **Lateness:** With the deadline set for the last day of classes, there is *no grace period* for late submission. If you are behind schedule, focus on having some things work and explain those things that don't. If everything is “99 percent okay” but nothing actually runs, that's a zero on the executable. Something that runs for the most part but has a couple of minor, documented bugs will get substantial partial credit.

Coursesite could slow right at the deadline. because of the many large submissions. That is why we ask that you submit a good draft version early. You need to plan ahead for these known delays and not submit at the last minute. Treat the deadline as if it were a few hours earlier, submit a draft, then download it and test it.

9. **Strict Policies:** It is very important that you take the time to ensure that the logistics of project submission are done right. If you make a simple goof (upload the wrong file or zip a wrong directory, etc.) it may seem unreasonable to give a grade of zero on the executable part considering all the work you did. However, lacking this policy, we would be making it trivial to gain an extension of time through an allegedly careless “goof”. While our course sizes are this large, we have to be strict in this regard. Check, recheck and be careful!

Note that computers seem to crash around project deadlines. You are responsible for backing up all your project code and data. While we take due precautions with our Oracle system, it is your responsibility to restore your database if it were to fail disastrously and all backups on our side were lost. We suggest backups to a personal external hard drive or USB drive.

Past experience suggests that the Oracle system will be slow leading up to the final project deadline. The slowness of the system will not be a reason for granting an extension of the deadline. If you are

working right up to the deadline, always have the latest version ready to submit just in case you don't get that one last test run done in time.

Collaboration:

- Your project database design and interface implementation is to be your own work with no outside help except from the instructional team.
- You may share raw data to load into your database or obtain data from public domain online data sources. Include a note in the README file as to the source of your data; also include a note if you have given data to someone else. Data sharing that is not documented in the README file constitutes an honor code violation.
- You may receive help on Unix/Linux logistics, but if that help is not from the instructional team, include a note in the README file.
- We intend to run project code through the Moss system to detect cases of possible plagiarism. In addition we may select some projects for "interview meetings" where you will discuss your projects with us sometime between the end of class and the end of finals.